

# Time Series Analysis & Forecasting Using R

An aerial photograph of a city skyline during sunset. The sky is a mix of orange, yellow, and light blue. In the foreground, there is a large green field, possibly a park or sports field, with some trees and a road. The city skyline is filled with various buildings, including several tall skyscrapers. The overall scene is a panoramic view of a modern city.

Transformation and decomposition

Bahman Rostami-Tabar

# Outline

- 1 Learning outcome
- 2 Per capita adjustments
- 3 Inflation adjustments
- 4 Mathematical transformations
- 5 Time series decompositions

# Outline

- 1 Learning outcome
- 2 Per capita adjustments
- 3 Inflation adjustments
- 4 Mathematical transformations
- 5 Time series decompositions

# Learning outcome

You should be able to:


- 1 Apply required adjustments on time series data before your analysis
- 2 Understand when time series decomposition would be helpful
- 3 Decompose time series using STL method

# Outline

- 1 Learning outcome
- 2 Per capita adjustments
- 3 Inflation adjustments
- 4 Mathematical transformations
- 5 Time series decompositions


# Per capita adjustments

```
global_economy |>  
  filter(Country == "Australia") |>  
  autoplot(GDP)
```



# Per capita adjustments

```
global_economy |>  
  filter(Country == "Australia") |>  
  autoplot(GDP / Population)
```



# Outline

- 1 Learning outcome
- 2 Per capita adjustments
- 3 Inflation adjustments**
- 4 Mathematical transformations
- 5 Time series decompositions




# Inflation adjustments

```
print_retail <- aus_retail |>
  filter(Industry == "Newspaper and book retailing") |>
  group_by(Industry) |>
  index_by(Year = year(Month)) |>
  summarise(Turnover = sum(Turnover))
aus_economy <- filter(global_economy, Code == "AUS")
print_retail |>
  left_join(aus_economy, by = "Year") |>
  mutate(Adj_turnover = Turnover / CPI) |>
  pivot_longer(c(Turnover, Adj_turnover),
    names_to = "Type", values_to = "Turnover"
  ) |>
  ggplot(aes(x = Year, y = Turnover)) +
  geom_line() +
  facet_grid(vars(Type), scales = "free_y") +
  labs(x = "Years", y = NULL,
    title = "Turnover: Australian print media industry")
```

# Inflation adjustments

Turnover: Australian print media industry



# Outline

- 1 Learning outcome
- 2 Per capita adjustments
- 3 Inflation adjustments
- 4 Mathematical transformations**
- 5 Time series decompositions

# Variance stabilization

If the data show different variation at different levels of the series, then a transformation can be useful.

# Variance stabilization

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as  $y_1, \dots, y_n$  and transformed observations as  $w_1, \dots, w_n$ .

# Variance stabilization

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as  $y_1, \dots, y_n$  and transformed observations as  $w_1, \dots, w_n$ .

## Mathematical transformations for stabilizing variation

Square root	$w_t = \sqrt{y_t}$	↓
Cube root	$w_t = \sqrt[3]{y_t}$	Increasing
Logarithm	$w_t = \log(y_t)$	strength

# Variance stabilization

If the data show different variation at different levels of the series, then a transformation can be useful.

Denote original observations as  $y_1, \dots, y_n$  and transformed observations as  $w_1, \dots, w_n$ .


## Mathematical transformations for stabilizing variation

Square root	$w_t = \sqrt{y_t}$	↓
Cube root	$w_t = \sqrt[3]{y_t}$	Increasing
Logarithm	$w_t = \log(y_t)$	strength

Logarithms, in particular, are useful because they are more interpretable: changes in a log value are **relative (percent) changes on the original scale**.

# Variance stabilization


```
food <- aus_retail |>  
  filter(Industry == "Food retailing") |>  
  summarise(Turnover = sum(Turnover))
```






# Variance stabilization

```
food |> autoplot(sqrt(Turnover)) +  
  labs(y = "Square root turnover")
```




# Variance stabilization

```
food |> autoplot(Turnover^(1 / 3)) +  
  labs(y = "Cube root turnover")
```




# Variance stabilization

```
food |> autoplot(log(Turnover)) +  
  labs(y = "Log turnover")
```



# Variance stabilization

```
food |> autoplot(-1 / Turnover) +  
  labs(y = "Inverse turnover")
```



# Box-Cox transformations

Each of these transformations is close to a member of the family of **Box-Cox transformations**:

$$w_t = \begin{cases} \log(y_t), & \lambda = 0; \\ (\text{sign}(y_t)|y_t|^\lambda - 1)/\lambda, & \lambda \neq 0. \end{cases}$$

# Box-Cox transformations

Each of these transformations is close to a member of the family of **Box-Cox transformations**:

$$w_t = \begin{cases} \log(y_t), & \lambda = 0; \\ (\text{sign}(y_t)|y_t|^\lambda - 1)/\lambda, & \lambda \neq 0. \end{cases}$$

- Actually the Bickel-Doksum transformation (allowing for  $y_t < 0$ )
- $\lambda = 1$ : (No substantive transformation)
- $\lambda = \frac{1}{2}$ : (Square root plus linear transformation)
- $\lambda = 0$ : (Natural logarithm)
- $\lambda = -1$ : (Inverse plus 1)

# Box-Cox transformations

# Box-Cox transformations

```
food |>  
  features(Turnover, features = guerrero)
```

```
## # A tibble: 1 x 1  
##   lambda_guerrero  
##             <dbl>  
## 1             0.0895
```



# Box-Cox transformations


```
food |>  
  features(Turnover, features = guerrero)
```

```
## # A tibble: 1 x 1  
##   lambda_guerrero  
##             <dbl>  
## 1             0.0895
```

- This attempts to balance the seasonal fluctuations and random variation across the series.
- Always check the results.
- A low value of  $\lambda$  can give extremely large prediction intervals.

# Box-Cox transformations

```
food |> autoplot(box_cox(Turnover, 0.0524)) +  
  labs(y = "Box-Cox transformed turnover")
```



# Transformations

- Often no transformation needed.
- Simple transformations are easier to explain and work well enough.
- Transformations can have very large effect on PI.
- If some data are zero or negative, then use  $\lambda > 0$ .
- `log1p()` can also be useful for data with zeros.
- Choosing logs is a simple way to force forecasts to be positive
- Transformations must be reversed to obtain forecasts on the original scale. (Handled automatically by `fable`.)

# Outline

- 1 Learning outcome
- 2 Per capita adjustments
- 3 Inflation adjustments
- 4 Mathematical transformations
- 5 Time series decompositions

# Time series decomposition

**Trend-Cycle** aperiodic changes in level over time.

**Seasonal** (almost) periodic changes in level due to seasonal factors (e.g., the quarter of the year, the month, or day of the week).

## Additive decomposition

$$y_t = S_t + T_t + R_t$$

where  $y_t$  = data at period  $t$

$T_t$  = trend-cycle component at period  $t$

$S_t$  = seasonal component at period  $t$

$R_t$  = remainder component at period  $t$

# STL decomposition

- STL: “Seasonal and Trend decomposition using Loess”
- Very versatile and robust.
- Seasonal component allowed to change over time, and rate of change controlled by user.
- Smoothness of trend-cycle also controlled by user.
- Optionally robust to outliers
- No trading day or calendar adjustments.
- Only additive.
- Take logs to get multiplicative decomposition.
- Use Box-Cox transformations to get other decompositions.


# US Retail Employment

```
us_retail_employment <- us_employment |>
  filter(year(Month) >= 1990, Title == "Retail Trade") |>
  select(-Series_ID)
us_retail_employment
```

```
## # A tsibble: 357 x 3 [1M]
##       Month Title      Employed
##       <mth> <chr>      <dbl>
## 1 1990 Jan Retail Trade 13256.
## 2 1990 Feb Retail Trade 12966.
## 3 1990 Mar Retail Trade 12938.
## 4 1990 Apr Retail Trade 13012.
## 5 1990 May Retail Trade 13108.
## 6 1990 Jun Retail Trade 13183.
## 7 1990 Jul Retail Trade 13170.
## 8 1990 Aug Retail Trade 13160.
## 9 1990 Sep Retail Trade 13113.
## 10 1990 Oct Retail Trade 13105.
```

# US Retail Employment

```
us_retail_employment |>  
  autoplot(Employed) +  
  labs(y = "Persons (thousands)", title = "Tot
```





# US Retail Employment

```
dcmp <- us_retail_employment |>  
  model(stl = STL(Employed))  
dcmp
```

```
## # A mable: 1 x 1  
##      stl  
##    <model>  
## 1    <STL>
```


# US Retail Employment

```
components(dcmp)
```

```
## # A dable: 357 x 7 [1M]
## # Key:      .model [1]
## # :        Employed = trend + season_year +
## #   remainder
##   .model      Month Employed  trend  season~1  remain~2
##   <chr>       <mth>    <dbl>   <dbl>    <dbl>    <dbl>
## 1 stl        1990 Jan    13256. 13288.   -33.0     0.836
## 2 stl        1990 Feb    12966. 13269.  -258.    -44.6
## 3 stl        1990 Mar    12938. 13250.  -290.    -22.1
## 4 stl        1990 Apr    13012. 13231.  -220.     1.05
## 5 stl        1990 May    13108. 13211.  -114.    11.3
## 6 stl        1990 Jun    13183. 13192.   -24.3    15.5
## 7 stl        1990 Jul    13170. 13172.   -23.2    21.6
## 8 stl        1990 Aug    13160. 13151.    -9.52   17.8
## 9 stl        1990 Sep    13113. 13131.  -39.5    22.0
## 10 stl       1990 Oct    13185. 13110.    61.6    13.2
```


# US Retail Employment

```
components(dcmp) |> autoplot()
```




# US Retail Employment

```
us_retail_employment |>  
  autoplot(Employed, color = "gray") +  
  autolayer(components(dcmp), trend, color = "  
  labs(y = "Persons (thousands)", title = "Tot
```



# US Retail Employment

```
components(dcmp) |> gg_subseries(season_year)
```



# Seasonal adjustment

- Useful by-product of decomposition: an easy way to calculate seasonally adjusted data.
- Additive decomposition: seasonally adjusted data given by


$$y_t - S_t = T_t + R_t$$

- Multiplicative decomposition: seasonally adjusted data given by

$$y_t/S_t = T_t \times R_t$$

# US Retail Employment

```
us_retail_employment |>  
  autoplot(Employed, color = "gray") +  
  autolayer(components(dcmp), season_adjust, color = "blue",  
    labs(y = "Persons (thousands)", title = "Total employment in US retail")
```



# Seasonal adjustment

- We use estimates of  $S$  based on past values to seasonally adjust a current value.
- Seasonally adjusted series reflect **remainders** as well as **trend**. Therefore they are not “smooth” and “downturns” or “upturns” can be misleading.
- It is better to use the trend-cycle component to look for turning points.



# STL decomposition


# STL decomposition

```
us_retail_employment |>  
  model(STL(Employed ~ trend(window = 15) + se  
    robust = TRUE  
  )) |>  
  components()
```

- `trend(window = ?)` controls wiggleness of trend component.
- `season(window = ?)` controls variation on seasonal component.
- `season(window = 'periodic')` is equivalent to an infinite window.

# STL decomposition


```
us_retail_employment |>  
  model(STL(Employed)) |>  
  components() |>  
  autoplot()
```



# STL decomposition

- `STL()` chooses season by default
- Can include transformations

```
us_retail_employment |>  
  model(STL(Employed)) |>  
  components() |>  
  autoplot()
```




# STL decomposition

- Algorithm that updates trend and seasonal components iteratively.
- Starts with  $\hat{T}_t = 0$
- Uses a mixture of loess and moving averages to successively refine the trend and seasonal estimates.
- trend window controls loess bandwidth on deasonalised values.
- season window controls loess bandwidth on detrended subseries.
- Robustness weights based on remainder.
- Default season: window = 13
- Default trend:

window = nextodd(ceiling((1.5\*period)/(1-


# Australian holidays

```
holidays |> autoplot(Trips) +  
  labs(y = "thousands of trips", x = "Year",  
       title = "Australian domestic holiday ni
```



# Australian holidays

```
holidays |>  
  model(stl = STL(Trips)) |>  
  components() |>  
  autoplot()
```



# Holidays decomposition


```
dcmp <- holidays |>
  model(stl = STL(Trips)) |>
  components()
dcmp
```

```
## # A dable: 640 x 8 [1Q]
## # Key:      State, .model [8]
## # :        Trips = trend + season_year +
## #   remainder
##   State .model Quarter Trips trend season_year
##   <chr> <chr>    <qtr> <dbl> <dbl>    <dbl>
## 1 ACT   stl     1998 Q1  196.  172.    -8.48
## 2 ACT   stl     1998 Q2  127.  157.     10.3
## 3 ACT   stl     1998 Q3  111.  142.   -16.8
## 4 ACT   stl     1998 Q4  170.  130.     14.6
## 5 ACT   stl     1999 Q1  108.  135.   -8.63
## 6 ACT   stl     1999 Q2  125.  148.     11.0
## 7 ACT   stl     1999 Q3  178.  166.   -16.0
## 8 ACT   stl     1999 Q4  218.  177.     13.2
```




# Holidays decomposition

```
dcmp |> gg_subseries(season_year)
```



# Holidays decomposition

```
autoplot(dcmp, trend, scale_bars = FALSE) +  
  autolayer(holidays, alpha = 0.4)
```



# Lab Session -transformation

## per capita

Consider the GDP information in `global_economy`. Plot the GDP per capita for each country over time. Which country has the highest GDP per capita? How has this changed over time?

## transformation

- 1 For the following series, find an appropriate transformation in order to stabilise the variance.
  - ▶ United States GDP from `global_economy`
  - ▶ Slaughter of Victorian “Bulls, bullocks and steers” in `aus_livestock`
  - ▶ Victorian Electricity Demand from `vic_elec`.

# Lab Session- decomposition

## 1 Produce the following decomposition

```
canadian_gas |>
  model(STL(Volume ~ season(window=7) + trend(window=11))) |>
  components() |>
  autoplot()
```

- 2 What happens as you change the values of the two window arguments?
- 3 How does the seasonal shape change over time? [Hint: Try plotting the seasonal component using `gg_season`.]
- 4 Can you produce a plausible seasonally adjusted series? [Hint: `season_adjust` is one of the variables returned by `STL`.]